

Process Management System

Computer Operating Systems
BLG 312E

2017-2018 Spring

Processes

- multiple jobs may be active at the same time
- each job may be a different running program

⇒ PROCESS

What is a Process?

Definition:

A process is a sequence of actions resulting from a run of a sequential program written for a specific function.

- process ⇔ task

What is a Process?

- process is a running program
- a process consists of a
 - sequential program code,
 - program counter,
 - register contents,
 - variables.

What is a Process?

- may be more than one process per program
- through system calls, processes
 - use system resources
 - communicate with each other
 - communicate with the world

Program ⇔ Process

Example: A programmer bakes a cake using a recipe.

recipe → program

ingredients → inputs

programmer → processor

Process → programmer reads recipe, obtains ingredients, performs necessary operations.

Program ⇔ Process

(*example cntd.*) His son enters the kitchen shouting that a bee has stung him. Programmer marks where he left off on the recipe, stops what he is doing, picks up the first aid book and starts the necessary treatment on his son.

treatment method → program
 medicines → inputs
 programmer → processor

Process → applying first aid using the treatment given in the book

Program ⇔ Process

(*example cntd.*)


Result: processor shared by two processes in time (time-sharing)

- the process which will have the processor is determined through an algorithm

Processes

- only one processor in system
 - time-sharing operation
 - "quantum"
 - only one of each system register
 - program counter, stack pointer, condition code register, general purpose registers, index register, ...
- ⇒ How is time-sharing achieved?

Time-Sharing




Time-Sharing

- the time a process will have the processor cannot be predicted
 - there should be no time dependent operations in a program code!

Time-Sharing

- process
 - RUNNING ⇒ has processor
 - NOT RUNNING ⇒ does not have processor




Question: Why would a process wait?

Process States


- processes are in different states throughout their lifetimes
- basic three states
 - running: using the processor
 - ready: can use the processor when it gets it
 - suspended: waits for an event; cannot use the processor even if it gets it

Process States




Process States

new process




Process States



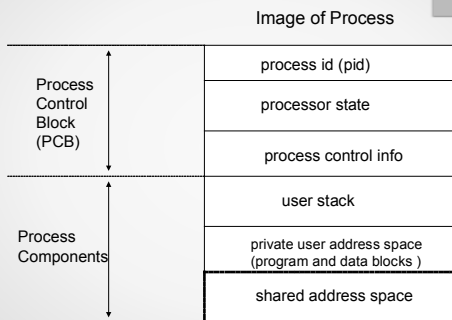
Implementing Processes

- operating system keeps info on all entities it manages
 - a different table for each
 - I/O tables
 - memory tables
 - file tables
 - proces tables**

Implementing Processes- Process Table



Implementing Processes – Image of a Process



Implementing Processes - PCB

- info regarding process in process descriptor field
 - process control block – PCB – holds info on process
- all operations on process through PCB
 - must have fast access to PCB
 - in some systems through a hardware register
 - in some systems special instructions to access PCB

Process Control Block (PCB) Contents

1. process identification info
 - process id
 - id of process' parent process
 - owner of process
2. current state of process and event it waits for (if any)
3. priority of process
4. scheduling info

Process Control Block (PCB) Contents

5. pointers to resources used by process
 - e.g. open files
6. pointer to virtual memory allocated to process
7. area where contents of system registers and user accessible processor registers are stored
 - general purpose registers, program counter, condition code register, index register, stack pointer,
 - ⇒ processor context


Operations on Processes

- create
- destroy
- suspend
- resume
- change priority

Creating Processes

- create process
 - in UNIX type systems only another process creates a process
 - a hierarchy exists among processes
 - creator process: **parent** proses
 - created process: **child** proses
 - a process may create multiple child processes

Process Hierarchy



Creating Processes

- when a process is created:
 - if process table full, process NOT created
 - if process table entry available,
 - process is assigned a unique id
 - process is assigned initial priorities
 - PCB is created and initialized
 - initial resources assigned (memory etc)
 - process is added to *ready* queue

Destroying Processes

- destroy a process
 - process removed from system
 - resources returned to system
 - pid returned to system
 - PCB and process table entry deleted
 - necessary operations on children performed
 - either keep entry until all children exit
 - or children assigned to another parent
 - e.g. in UNIX to the *init* process (pid=1)

Other Operations on Processes

- suspend a process
 - for short term suspension, resources are not removed
 - for long term suspension, depending on resources, some may be removed
- resume a process
 - to resume operation of a process from where it left off in its previous execution
- change priority of a process

Steps during a Process State Change

- processor context saved
- PCB of running process updated
- running process added to appropriate queue (ready / suspended)
- new process to run determined
- PCB of selected process updated
- info on memory management updated
- context of selected process loaded onto registers

